

Tracing Machine-Advice Markers Across OLMo-2: Open Training Data, Stagewise Dynamics, and a Measurement Audit

Anonymous ACL submission

Abstract

Instruction-tuned models often combine empathy, validation, professional referral, and structuring language when giving personal advice. We ask whether this is one template and where its targeting changes during training. We trace these markers across OLMo-2’s public Base, SFT, DPO, and RLVR checkpoints using a 140-prompt behavioral battery (five conditions; five generations per prompt), with its 90 advice-plus-factual prompts used for stagewise comparisons, and exact searches of the corresponding open training data. The pattern decomposes: validation is the only marker elevated on advice against every matched control; empathy tracks emotional content, structuring tracks how-to form, and referral also fires on clinical-information prompts. Across two model sizes, all eight advice-versus-factual gaps widen at SFT, with six surviving multiplicity correction. DPO further widens expressive-marker gaps but spreads structuring to factual answers, collapsing its advice gap; no later transition is significant. Exact data provide compatible phrase-level evidence: referral phrases occur in SFT assistant turns, referral and structuring phrases are enriched in DPO chosen responses under paired, word-normalized tests, and both are nearly absent from the searched RLVR data. A separate 120-sentence AI consensus audit finds detector macro precision 0.28 and recall 0.72, so checkpoint-level rates remain provisional. Because SFT evidence is phrase-sparse, the strongest conclusion is the paired, length-normalized DPO enrichment rather than semantic behavior prevalence across checkpoints.

1 Introduction

Ask a modern instruction-tuned model “I’ve been feeling really anxious lately and can’t sleep, any advice?” and you will, in our sample, reliably receive a structured response: an acknowledgment, validation of the user’s emotional state, suggestion to seek professional help, and list-organizing lan-

guage. Two questions follow. First, is this a single *advice template*, or several behaviors that merely co-occur on advice prompts? Second, *where* in training does each behavior emerge?

Answering either requires public training artifacts and intermediate weights. A final OLMo-2 Instruct checkpoint is the product of four data-bearing stages: pretraining, supervised fine-tuning (SFT), direct preference optimization (DPO), and reinforcement learning with verifiable rewards (RLVR) (Team OLMo et al., 2024). OLMo-2 is, to our knowledge, the only current family for which all post-training data *and* intermediate weight checkpoints are public. We exploit both: we run the same behavioral evaluation on every public stage checkpoint, and we search each stage’s exact data: the Stage-1 corpus OLMo-Mix-1124 via infini-gram, each model’s exact SFT mixture (tulu-3-sft-olmo-2-mixture for the 1124-7B line, tulu-3-sft-olmo-2-mixture-0225 for the 0425-1B line), the DPO preference mixes, and the RLVR sets.

We are deliberate about what this evidence is: **observational co-occurrence and stage localization, not causal attribution**. We measure where behaviors appear across checkpoints and where their phrasings live in the corresponding data; we do not retrain. Our data search is not fully end-to-end: Stage-2 mid-training (Dolmino-Mix-1124) has no public index and is not searched.

Contributions.

1. **An auditable behavior-level tracer on exact open data.** A 140-prompt battery in five matched conditions (advice, factual, emotional non-advice, neutral advice, domain factual), five generations per prompt, a released lexicon+e5 detector, role-scoped corpus search, and a separate 120-sentence AI consensus audit.

2. **A controlled decomposition of the “advice template.”** At the primary threshold, detector-positive empathy tracks emotional content, structuring language tracks the how-to form, referral markers track advice/health, and the validation marker alone is elevated against every matched control; the threshold sweep identifies which distinctions are detector-sensitive (Section 4.1).
3. **A stagewise difference-in-differences with significance.** A prompt-clustered permutation test on per-prompt trajectories across Base→SFT→DPO→RLVR checkpoints shows every detector gap widens at SFT (six of eight changes survive multiplicity correction), DPO significantly widens the expressive gaps and *collapses* the structuring gap (generalizing lists to non-advice) but not the near-ceiling referral gap, with no later-stage gap change detected (Section 4.2).
4. **A measurement audit rather than an assumed oracle.** Two independent blinded AI passes agree exactly on 117/120 sentences; Claude Opus 4.8 adjudicates the remaining three. The resulting low precision sharply limits output-level interpretation, while the paired DPO and exact corpus results remain detector-independent at the phrase level.
5. **Compatible data-layer evidence with paired inference.** Referral phrasings cluster in exact-SFT assistant turns (phrase-sparse, leave-one-out reported) and remain enriched in DPO chosen responses under pair-level McNemar tests and word normalization; searched RLVR counts are near zero. Verbatim-only novelty is stated as a first-class limitation (Sections 4.3–4.4).

2 Related Work

Behavioral tests and checkpoint trajectories. Our matched prompt battery follows the behavioral-testing principle that targeted capabilities require controlled tests beyond aggregate accuracy (Ribeiro et al., 2020). Pythia makes dense public checkpoints available for training-dynamics research (Biderman et al., 2023); checkpoint studies have tracked how loss and downstream behavior change during pretraining (Xia et al., 2023). We instead compare named post-training stages and test changes in condition gaps rather than raw behavior

rates.

Pretraining attribution and n-gram search. infini-gram (Liu et al., 2024) makes exact-substring counting over web-scale corpora tractable, and OLMoTrace (Liu et al., 2025) traces model outputs back to pretraining text. Both are pretraining-only; we add role-scoped search over the post-training layers, where our stagewise results locate the change.

Open instruction and preference data. Tulu-3 (Lambert et al., 2024) is the recipe behind OLMo-2’s post-training; the SFT mixtures, DPO preference mixes, and RLVR sets are public, as is oasst1 (Köpf et al., 2023) (itself an internal subset of the Tulu-3 mixtures). Public data at every stage is what makes the trace auditable.

The superficial alignment hypothesis. LIMA (Zhou et al., 2023) argues alignment largely surfaces a latent style from few demonstrations; Ghosh et al. (2024) interrogate this. Our stagewise gaps give a behavior-level version: each marker is nonzero at base, and post-training produces larger advice-minus-factual detector gaps.

Causal data attribution. LESS uses optimizer-aware gradient similarity to select influential instruction data (Xia et al., 2024); Xiao and Aranguri (2026) perform probe-based causal attribution on OLMo-2 DPO and mitigate a behavior by editing top-attributed preference pairs. We differ in kind: our trace is observational and cheap, covers all four stages and both data and checkpoints, and *nominates* clusters and stages for exactly such interventions; it does not establish causation.

3 Method

3.1 Behavior detection (lexicon + e5)

A hand-seeded, inspectable lexicon defines four reported constructs: **empathy expression** (artifact key `empathy_opener`), **validation**, **professional referral** (`disclaimer`), and **structuring language** (`structure`). A fifth category, `crisis_referral`, fails the 120-sentence audit and is excluded. Exact matching is augmented with **e5-small** paraphrase embeddings (Wang et al., 2024), a response is detector-positive when any sentence exceeds similarity $\tau = 0.82$ to a seed exemplar (sweep 0.80/0.82/0.84). **Detector construction and circularity.** The lexicon, exemplars, and threshold were hand-seeded during an 8-prompt

pilot on OLMo-2-1B and frozen *before* the present 140-prompt battery was authored; all three control conditions and 100 of the 140 prompts were written after freezing. The sample sizes refer to distinct datasets: the 140 prompts form the behavioral evaluation, whereas 120 stratified final-model sentences form a detector-accuracy audit. We audit those sentences (40/30/20/20/10 across the five conditions) using two complete, independent, blinded AI passes and Claude Opus 4.8 adjudication of their three disagreements. This is an AI diagnostic, not human validation. The labels, adjudication reasons, and detector are released (Section 4.4; Appendix B).

3.2 Prompt battery and generation

Five conditions isolate the axes the “advice template” conflates: **advice** (60 personal/emotional help requests), **factual** (30 knowledge questions), **emotional non-advice** (20 descriptive prompts about emotional topics, no help request), **neutral advice** (20 how-to requests with no emotional stakes), and **domain factual** (10 impersonal clinical-information questions). Each prompt is answered with **five sampled generations** (temperature 0.7, top- p 0.9, 256 new tokens, distinct seeds).

3.3 Statistical unit

The clustering unit is the **prompt**: each prompt contributes its mean detector-positive indicator over five generations, and condition rates are means over prompts. All between-condition inference uses **prompt-clustered permutation tests** (10,000 permutations of prompt-level means; two-sided; Monte Carlo $p = (b + 1)/(B + 1)$). Benjamini–Hochberg correction is applied separately to the 24 matched-control, 28 adjacent-stage, and 8 DPO pairwise families. We do *not* treat generations as independent Bernoulli trials: Fisher/Wilson statistics computed on “equivalent counts” are reported only as an appendix sensitivity analysis with that caveat (Appendix D).

3.4 Stagewise evaluation and stage-data search

The same evaluation runs on every public checkpoint: 1B Base→SFT→DPO→RLVR1→Instruct; 7B Base→SFT→DPO→Instruct. Base checkpoints have no chat template and use a plain Question/Answer format (a comparison caveat we state). The stagewise grid runs the 90 advice+factual prompts; all five conditions run on the final models.

For each stage’s data we search: the exact Stage-1 corpus (v4_olmo-mix-1124_llama; novelty = 1−verbatim coverage), each model’s exact SFT mixture in DuckDB with **assistant-only** scoping (866,138 / 939,344 conversations for 1B/7B; oasst1 not searched separately since it is an internal subset), the DPO preference mixes (pair-level chosen-only/rejected-only counts, exact McNemar tests, and phrase occurrences per million response words), and the RLVR sets. SFT recoverability compares behavioral phrases against **length-matched topical controls** from the same answers, per-million documents, with a distinct-phrase leave-one-out.

Reproducibility. Stagewise, condition, and corpus numbers derive from `summary_v3.json` and released scripts under `experiments/`; prompts, lexicon, seeds, thresholds, and adjusted- p output ship with it. Audit numbers are separately materialized in `ai_consensus_validation_v3.json`, with both 120-item passes, consensus labels, Opus decisions, and reasons released alongside it.

4 Results

Results at a glance. Four outcomes organize the analysis. First, the apparent advice bundle decomposes into affect-linked empathy, form-linked structuring, domain-linked referral, and a validation marker consistently elevated on advice. Second, all eight advice-minus-factual detector gaps widen at SFT (six survive correction), while DPO selectively widens expressive gaps and spreads structuring to factual answers. Third, exact phrase searches find referral and structuring enrichment in DPO chosen responses but almost none in the searched RLVR data. Fourth, the separate 120-sentence audit finds low detector precision, so the checkpoint results are marker trajectories rather than estimates of semantic behavior prevalence.

4.1 The detector-positive advice bundle decomposes at the primary threshold

Table 1 gives mean prompt-level detector-positive rates per condition on the final models; Table 2 tests advice against each matched control. Against *plain factual* controls every detector category separates sharply in both models (all prompt-clustered $p < 10^{-4}$; referral 0.927 vs 0.073 at 1B, 0.913 vs 0.060 at 7B). But the matched controls show this is **not one advice-specific template**:

- **Empathy expression tracks emotional content at $\tau = .82$.** It fires *more* on emotional non-advice descriptions than on advice (0.79 vs 0.587 at 1B, $p = .005$; 0.76 vs 0.617 at 7B, $p = .098$), and near-zero on neutral advice (0.03–0.07).
- **Structuring language tracks the how-to form.** It is at ceiling on *neutral* advice exactly as on personal advice (diff 0.00, $p = 1.0$ at 1B; -0.03 , $p = .57$ at 7B).
- **Professional referral tracks the advice/health domain at $\tau = .82$.** It fires on 0.66–0.72 of neutral-advice and 0.80–0.82 of clinical-information prompts; its advice increment over domain-factual is small and not significant ($+0.127$, $p = .053$ at 1B; $+0.093$, $p = .159$ at 7B).
- **The validation marker is the only detector category consistently elevated on personal advice relative to every matched control:** higher on advice than *every* control in *both* models (all unadjusted $p \leq .02$; all BH-adjusted $q \leq .028$), including emotional non-advice (0.36 vs 0.16, $p = .008$; 0.383 vs 0.13, $p = .003$).

Within the frozen detector, the primary-threshold marker bundle separates into affect-linked empathy, form-linked structuring, domain-linked referral, and a validation marker consistently elevated on advice. These are not detector-invariant universals: the threshold sweep preserves the structuring result but qualifies the empathy and referral contrasts (Section 4.4). The 1B-vs-7B differences are small and mixed; we make no scaling claim.

4.2 Stagewise marker dynamics: every gap widens at SFT

Raw advice-side rates rising across stages would be ambiguous because post-training could simply make *all* outputs more cautious and listy. The right quantity is the **advice-minus-factual gap** Δ_s at each stage (Table 4) and its *change* across adjacent stages, a stage $\times$ condition difference-in-differences tested with a prompt-clustered permutation test on per-prompt stage trajectories (Table 3; 10k permutations, $n=60$ advice/30 factual). Three findings.

- **SFT widens every gap.** All eight base $\rightarrow$ SFT changes are positive (raw $p < 0.05$; referral $+0.33/+0.28$, structuring $+0.39/+0.23$).

BH correction across all 28 adjacent-stage $\times$ behavior tests preserves six; validation is marginal ($q = .087/.055$). Every detector category is already nonzero at base, so SFT widens advice-minus-factual marker gaps for categories already detected at Base.

- **DPO acts marker-specifically.** It significantly further widens the *expressive-marker* gaps (validation $+0.16/+0.15$, $p \leq 0.008$; 1B empathy $+0.15$, $p = 0.005$) but *not* the referral gap ($+0.10/+0.08$, $p = 0.067/0.109$), which is already near ceiling on advice after SFT. Its preference data nonetheless contains more referral phrasings in chosen than rejected responses (Table 5), consistent with *maintaining* that near-saturated gap.
- **DPO collapses the structuring gap.** Its gap falls $-0.38/-0.28$ at DPO ($p < 10^{-4}$), not because advice-side structuring drops (it is at ceiling) but because DPO raises structuring on *factual* answers from 0.07 to 0.54–0.59. Preference tuning spreads structuring language everywhere while leaving the other advice-minus-factual detector gaps positive.

RLVR and the final merge move no gap significantly (every dpo $\rightarrow$ rlvr and final-merge transition has $p \geq 0.21$), consistent with RLVR’s math-only data (Section 4.3). These significance tests are internally valid for the frozen detector outputs; because the audit finds condition-dependent false positives and does not span checkpoints, they do not by themselves establish changes in human-judged behavior prevalence.

4.3 The data layers provide compatible phrase-level evidence

SFT: a finite, auditable referral cluster (phrase-sparse). Assistant-only search of each model’s exact mixture surfaces thousands of referral demonstrations: “healthcare professional” (3075/3106 for 1B/7B), “mental health professional” (2109/2120), “professional help” (1765/1774), “seek professional” (1353/1355). Length-matched recoverability is above topical controls at ≥ 3 words (3-word: $19.5\times$ at 1B, $7.3\times$ at 7B; 4-word: $141.7\times/14.5\times$) but **phrase-sparse**: only 5/6 distinct 3-word types, dominated by the list preamble “here are some.” The distinct-phrase leave-one-out is the honest figure: dropping that one phrase lowers the 3-word mean from 3358.1 to **758.8/M** (1B) and 3026.0

Behavior	Model	Advice	Factual	Emotional	Neutral adv.	Domain fact.
empathy expr.	1B	0.587	0.033	0.790	0.030	0.340
	7B	0.617	0.027	0.760	0.070	0.360
validation	1B	0.360	0.000	0.160	0.000	0.120
	7B	0.383	0.013	0.130	0.030	0.060
referral	1B	0.927	0.073	0.510	0.660	0.800
	7B	0.913	0.060	0.540	0.720	0.820
structuring	1B	1.000	0.587	0.820	1.000	0.940
	7B	0.970	0.547	0.830	1.000	0.840

Table 1: Mean prompt-level detector-positive rate per condition (final models, $\tau = 0.82$; each prompt averaged over 5 generations). Bold marks the cells that break the “advice-specific” reading: empathy peaks on emotional *non-advice* prompts, and structuring language is at ceiling on *neutral* advice. Rates are upward-biased (Section 4.4); *crisis_referral* is excluded (0 TP / 11 FP in the 120-sentence audit).

Behavior	M	−Emot.	−Neut.adv.	−Dom.fact.
empathy	1B	−.203 (.005)	.557 ($<10^{-4}$)	.247 (.010)
	7B	−.143 (.098)	.547 ($<10^{-4}$)	.257 (.031)
validation	1B	.200 (.008)	.360 ($<10^{-4}$)	.240 (.020)
	7B	.253 (.003)	.353 (.0001)	.323 (.004)
referral	1B	.417 ($<10^{-4}$)	.267 ($<10^{-4}$)	.127 (.053)
	7B	.373 ($<10^{-4}$)	.193 (.001)	.093 (.159)
structuring	1B	.180 ($<10^{-4}$)	.000 (1.0)	.060 (.017)
	7B	.140 (.004)	−.030 (.574)	.130 (.026)

Table 2: Advice minus each matched control (prompt-clustered permutation p , 10k draws). Benjamini–Hochberg correction across the 24 displayed tests leaves the validation conclusion unchanged (all $q \leq .028$); non-significant and negative cells carry the decomposition (Section 4.1).

Model	Behavior	base→SFT	SFT→DPO
1B	empathy	+0.197 ($<10^{-3}$)	+0.153 (.005)
1B	validation	+0.110 (.040)	+0.157 (.007)
1B	referral	+0.333 ($<10^{-4}$)	+0.097 (.067)
1B	structuring	+0.393 ($<10^{-4}$)	−0.377 ($<10^{-4}$)
7B	empathy	+0.277 ($<10^{-3}$)	+0.070 (.251)
7B	validation	+0.120 (.024)	+0.147 (.008)
7B	referral	+0.277 ($<10^{-4}$)	+0.083 (.109)
7B	structuring	+0.233 ($<10^{-3}$)	−0.277 ($<10^{-4}$)

Table 3: Change in the advice-minus-factual gap across transitions (plus-one-corrected prompt-clustered permutation p , 10k draws). BH correction across all 28 adjacent-stage×behavior tests preserves six of eight SFT effects (validation: $q = .087/.055$) and all stated DPO effects; later row $p \geq .21$.

to **1065.9/M** (7B) (Appendix E). Any inference must use the 5–6 distinct types as its unit, not the 46/55 occurrences. Role-scoping remains necessary: blob search inflates “talk to a” 1.41–1.43× via user turns.

DPO: paired and length-normalized enrichment.

The exact mixes contain 378,301/378,341 preference pairs. Chosen responses are slightly *shorter* in aggregate (chosen/rejected mean-word ratio 0.990/0.997), ruling out a general length explanation. At pair level, professional-referral language occurs in 2478 vs 2098 chosen-only/rejected-only pairs at 1B and 2602 vs 1925 at 7B (exact McNemar, BH $q = 4.1 \times 10^{-8} / 2.1 \times 10^{-23}$); per-million-word enrichment remains 1.12×/1.25×. Structuring language is stronger still (27047 vs 21779 and 28466 vs 19838; $q < 10^{-124}$; word-normalized 1.16×/1.27×). Empathy is weakly enriched only at 1B after correction, and validation is not significant (Table 5). These paired co-occurrences are compatible with, but do not attribute or directly validate, the checkpoint changes.

RLVR: a clean negative. The RLVR datasets (GSM/MATH/IF) contain essentially none of the advice phrasings: every referral, empathy, and validation phrase counts 0–1 across RLVR–GSM, RLVR–MATH, and the IF-constraints set (only generic markers appear: “first,” 881, “here are some” 48 in the IF set). We find no direct lexical evidence that the searched RLVR sets introduce these behaviors. Their stability through the RLVR stage (Table 4) is consistent with retention rather than new introduction at RLVR.

Pretraining: not a verbatim echo (secondary, hedged). Against the exact Stage-1 corpus, advice answers are more novel than factual ones (0.559 vs 0.383 at 1B; 0.525 vs 0.379 at 7B). This is verbatim-only, length-confounded, and whole-answer rather than behavior-span, so we read it only as: advice answers are not verbatim OLMo-Mix copies. The base checkpoints’ nonzero behavior rates are the stronger evidence that the *capability* predates post-training.

Model	Stage	Empathy	Validation	Referral	Structuring
1B	base	.207 (.194)	.073 (.073)	.473 (.413)	.730 (.397)
1B	sft	.390 (.390)	.183 (.183)	.753 (.746)	.857 (.790)
1B	dpo	.590 (.543)	.353 (.340)	.917 (.844)	1.000 (.413)
1B	rlvr	.597 (.550)	.373 (.373)	.903 (.816)	1.000 (.400)
1B	instruct	.587 (.554)	.360 (.360)	.927 (.854)	1.000 (.413)
7B	base	.217 (.217)	.070 (.063)	.543 (.490)	.847 (.480)
7B	sft	.493 (.493)	.183 (.183)	.767 (.767)	.787 (.714)
7B	dpo	.597 (.564)	.330 (.330)	.923 (.850)	.977 (.437)
7B	instruct	.617 (.590)	.383 (.370)	.913 (.853)	.970 (.423)

Table 4: Stagewise advice detector-positive rate with the **advice-minus-factual gap in parentheses** ($\tau = 0.82$). The referral gap widens at SFT (+0.33/+0.28, $p < 10^{-4}$; interaction tests in Table 3) and is not significantly changed thereafter; the structuring gap collapses at DPO because factual answers become listy too. Base checkpoints use a plain QA format (no chat template).

M	Behavior	C-only	R-only	word C/R	BH q
1B	empathy	1768	1634	1.04	.036
1B	validation	183	162	1.12	.282
1B	referral	2478	2098	1.12	$< 10^{-7}$
1B	structuring	27047	21779	1.16	$< 10^{-124}$
7B	empathy	1734	1622	1.05	.074
7B	validation	191	156	1.23	.078
7B	referral	2602	1925	1.25	$< 10^{-22}$
7B	structuring	28466	19838	1.27	$< 10^{-300}$

Table 5: DPO pair-level and length-normalized analysis. C/R is the ratio of phrase occurrences per million whitespace-delimited response words; q is Benjamini–Hochberg adjusted across all eight exact McNemar tests.

4.4 Detector audit and threshold sensitivity

A condition-stratified random audit samples 120 final-model sentences across all five conditions (40 advice, 30 factual, 20 emotional non-advice, 20 neutral advice, 10 domain factual). The sample contains 56 1B and 64 7B sentences from 97 distinct seed-0 responses; it was not stratified on detector predictions, and multiple sentences can come from one response. Two independent blinded Codex passes agree exactly on 117/120 sentences (97.5%); category agreement is 99.2–100%, with Cohen’s $\kappa = .885/.663/1.000/.978$ for empathy/validation/referral/structuring. Claude Opus 4.8 adjudicates the three disagreements. This establishes reproducible AI consensus, not human ground truth.

Against that consensus, the detector has **low precision and category-dependent recall** (Table 6; four-category macro P 0.28, R 0.72, F1 0.36; empathy and validation recall are each 0.50). The errors are also condition-dependent: professional referral produces 12 false positives on advice, one on emotional non-advice, two on neutral advice,

and none on factual or domain-factual items; validation’s two false positives are both advice items. Thus subtraction does not guarantee cancellation, and even cross-condition directions remain provisional. `crisis_referral` again fails (0 TP / 11 FP) and is excluded. Absolute output rates are therefore upward-biased (not formal upper bounds because recall is below one), while the stagewise results should be read as trajectories of a frozen marker detector, not estimates of semantic behavior prevalence.

The low positive prevalence yields wide Wilson 95% intervals (e.g., referral precision [.011, .283] and validation precision [.061, .792]; all intervals ship with the artifact), and the sentence-level sample does not directly validate answer-level emission. The audit itself is AI-only. We therefore also release a blinded two-human workflow with separate files and adjudication, but do not claim it was completed. The threshold sweep is informative rather than uniformly reassuring (Appendix C): structuring’s advice–neutral contrast stays within .07 in all six model-threshold cells and validation exceeds every control in five of six, but empathy reverses at .84 and referral changes magnitude. These are robustness checks on the detector, not a validated taxonomy.

5 Discussion

From “advice template” to trained marker trajectories. The detector-positive decomposition and stagewise gaps nominate a testable account: relevant language is already detected at Base; SFT widens every advice-minus-factual marker gap (six of eight changes survive correction); DPO then acts selectively by significantly widening the expressive gaps and collapsing the structuring gap

category	P	R	F1	tp/fp/fn
empathy expr.	0.22	0.50	0.31	2/7/2
validation	0.33	0.50	0.40	1/2/1
referral	0.06	1.00	0.12	1/15/0
structuring	0.49	0.87	0.63	26/27/4
macro P 0.28 / R 0.72 / F1 0.36; $n=120$				

Table 6: Detector P/R/F1 against the adjudicated 120-sentence AI consensus ($\tau = 0.82$). The two blinded AI passes agreed on 117/120 items; Opus adjudicated three. This is not human validation. Wilson intervals are released with the artifact. `crisis_referral`: 0 TP / 11 FP, dropped.

(spreading organizers to factual answers) while leaving the already-saturated referral gap unmoved, even as referral phrasings show strong paired, word-normalized chosen-response enrichment; no RLVR-stage gap change is detected. The output trajectory remains provisional because of measurement error. The strongest evidence is instead phrase-level and independently auditable: the DPO chosen/rejected differential and RLVR zero-counts do not rely on the embedding threshold, although they still establish co-occurrence rather than causation.

Safety implications. Professional-referral phrasing, the bundle’s safety-relevant marker, localizes to a finite, inspectable SFT demonstration cluster and measurable DPO chosen-response enrichment. That makes the marker auditable, and also fragile: mixture or preference-annotation changes could silently weaken it. The causal leave-cluster-out test our trace nominates (remove the referral cluster, re-SFT, compare against random-cluster removal) is the direct check.

Methodological lessons. (1) Occurrence-weighted provenance ratios are fragile: report distinct phrase types and a leave-one-out. (2) Advice-side rates across stages are ambiguous; report stage $\times$ condition *gaps*. (3) Whole-conversation search miscredits user turns; scope to assistant turns. (4) Cluster inference at the prompt, not the sample.

6 Limitations

- **Observational throughout.** Stage localization and data co-occurrence, including the DPO chosen/rejected differential, establish no causation; the nominated leave-cluster-out re-SFT is future work.

- **Detector and threshold.** The complete 120-sentence AI consensus audit finds low precision and category-dependent recall (referral P 0.06; empathy and validation R 0.50) plus condition-dependent errors. The two blinded passes agree strongly and Opus adjudicates all three disagreements, but this is not human validation. Rates are upward-biased; absolute levels, cross-condition gaps, and checkpoint trajectories remain provisional. Empathy/referral are threshold-sensitive; structuring and validation are more stable under the threshold sweep but still show false positives. A provenance-safe two-human workflow is released.
- **Coverage gaps.** Stage-2 mid-training (Dolmino-Mix-1124) is unsearched (no public index); matched controls exist only at the final stage; base checkpoints use a different (plain QA) prompt format.
- **Phrase sparsity.** The SFT co-occurrence rests on 5–6 distinct 3-word types dominated by one list marker; use the leave-one-out figures.
- **Statistical scope.** We apply Benjamini–Hochberg correction across the full 28-test adjacent-stage $\times$ behavior grid and, separately, across the 24 matched-control tests. Six of eight Base $\rightarrow$ SFT changes survive; the two validation changes are marginal. Five seeds bound but do not eliminate sampling variance; prompt counts are modest; one model family, two sizes, one decoding configuration; no scaling claim.
- **Novelty** is verbatim-only, length-confounded, and whole-answer.

7 Conclusion

Machine-advice language in OLMo-2 is not one indivisible template. A frozen detector yields distinct affect-, form-, and domain-linked trajectories: all eight advice–factual marker gaps widen at SFT (six survive multiplicity correction), while DPO widens expressive-marker gaps and spreads structuring markers to factual answers. The 120-sentence audit, however, finds enough condition-dependent false positives that these output trajectories must remain provisional. The strongest result is detector-independent at the matching stage: exact professional-referral and structuring phrases are enriched in DPO chosen responses under paired tests

and word normalization, while the searched RLVR data contain almost none. The paper’s contribution is therefore both a trace and an audit of where that trace is trustworthy. We release the tracer, battery, AI consensus labels, Opus decisions, and machine-readable results. The nominated next steps are causal and measurement-side: leave-cluster-out re-SFT against random-cluster controls (cf. [Xiao and Aranguri, 2026](#)), completion of the stratified two-annotator validation, and a self-hosted Stage-2 index.

Ethics Statement

Data and human subjects. This work uses only publicly released artifacts (OLMo-Mix-1124 via the public infinigram API; the OLMo-2 Tulu-3 SFT mixtures, DPO preference mixes, and RLVR sets; open-weight OLMo-2 models and their public intermediate checkpoints). We collected no new data from people. All 140 prompts are short, author-written strings describing hypothetical situations with no personal information. The study involves no human subjects and needed no ethics-board review.

Sensitive domain. Our prompts touch emotional-support, health, and relationship topics, and our analysis concerns safety-relevant moves such as professional referrals. Given the detector’s low precision we report detector-positive rates as upward-biased and drop `crisis_referral`. Nothing here should be read as evidence that the studied models reliably provide crisis support; they were not evaluated for clinical accuracy and are not mental-health tools.

Dual use. Making the provenance of a safety behavior inspectable could in principle guide its removal from a training mixture; we frame this as a risk to monitor, show co-occurrence rather than causation, and release an analysis tool, not any model or dataset that could cause direct harm.

Use of AI Assistants

AI coding and writing assistants were used to implement the analysis tooling and to draft and edit the manuscript. Two independent blinded Codex passes annotated the 120-sentence detector audit; Claude Opus 4.8 adjudicated their three disagreements. We report this provenance explicitly and do not describe the audit as human validation. The author designed the study, chose the methodology and

metrics, and made the scientific decisions, including dropping `crisis_referral` and choosing the statistical unit. All reported numbers were computed by released code from released artifacts and checked against them.

References

- Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, Usven Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar Van Der Wal. 2023. [Pythia: A suite for analyzing large language models across training and scaling](#). In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 2397–2430. PMLR.
- Sreyan Ghosh, Chandra Kiran Reddy Evuru, Sonal Kumar, S. Ramaneswaran, Deepali Aneja, Zeyu Jin, Ramani Duraiswami, and Dinesh Manocha. 2024. [A closer look at the limitations of instruction tuning](#). In *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- Andreas Köpf, Yannic Kilcher, Dimitri von Rütte, Sotiris Anagnostidis, Zhi-Rui Tam, Keith Stevens, Abdullah Barhoum, Nguyen Minh Duc, Oliver Stanley, Richárd Nagyfi, Shahul ES, Sameer Suri, David Glushkov, Arnab Dantuluri, Andrew Maguire, Christoph Schuhmann, Huu Nguyen, and Alexander Mattick. 2023. [OpenAssistant conversations – democratizing large language model alignment](#). In *Advances in Neural Information Processing Systems 36 (NeurIPS), Datasets and Benchmarks Track*.
- Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V. Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, Yuling Gu, Saumya Malik, Victoria Graf, Jena D. Hwang, Jiangjiang Yang, Ronan Le Bras, Øyvind Tafjord, Chris Wilhelm, Luca Soldaini, and 4 others. 2024. [Tulu 3: Pushing frontiers in open language model post-training](#). Preprint, arXiv:2411.15124.
- Jiacheng Liu, Taylor Blanton, Yanai Elazar, Sewon Min, Yen-Sung Chen, Arnavi Chheda-Kothary, Huy Tran, Byron Bischoff, Eric Marsh, Michael Schmitz, Cassidy Trier, Aaron Sarnat, Jenna James, Jon Borchardt, Bailey Kuehl, Evie Yu-Yen Cheng, Karen Farley, Taira Anderson, David Albright, and 11 others. 2025. [OLMoTrace: Tracing language model outputs back to trillions of training tokens](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL): System Demonstrations*, pages 178–188, Vienna, Austria. Association for Computational Linguistics.
- Jiacheng Liu, Sewon Min, Luke Zettlemoyer, Yejin Choi, and Hannaneh Hajishirzi. 2024. [Infinigram: Scaling unbounded \$n\$ -gram language models to a](#)

666	trillion tokens. In <i>First Conference on Language Modeling (COLM)</i> .	Examples.	718
667			
668	Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin,	• Advice: I've been feeling really	719
669	and Sameer Singh. 2020. <i>Beyond accuracy: Be-</i>	anxious lately and can't	720
670	<i>havioral testing of NLP models with CheckList</i> . In	sleep. Any advice?	721
671	<i>Proceedings of the 58th Annual Meeting of the Asso-</i>	• Factual: How does photosynthesis	722
672	<i>ciation for Computational Linguistics</i> , pages 4902–	work?	723
673	4912. Association for Computational Linguistics.		
674	Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groen-	• Emotional non-advice: Describe the	724
675	eveld, Kyle Lo, Shane Arora, Akshita Bhagia, Yuling	common symptoms of loneliness.	725
676	Gu, Shengyi Huang, Matt Jordan, Nathan Lambert,	• Neutral advice: How should I	726
677	Dustin Schwenk, Oyvind Taffjord, Taira Anderson,	organize my bookshelf?	727
678	David Atkinson, Faeze Brahman, Christopher Clark,	• Domain factual: What are common	728
679	Pradeep Dasigi, Nouha Dziri, and 23 others. 2024.	clinical treatments for	729
680	<i>2 OLMo 2 Furious</i> . Preprint, arXiv:2501.00656. A	insomnia?	730
681	shorter version appeared at COLM 2025.		
682	Liang Wang, Nan Yang, Xiaolong Huang, Linjun Yang,	B Behavior lexicon and detector	731
683	Rangan Majumder, and Furu Wei. 2024. <i>Multilingual</i>	threshold	732
684	<i>e5 text embeddings: A technical report</i> . Preprint,		
685	arXiv:2402.05672.	A category fires for an answer if <i>either</i>	733
686	Mengzhou Xia, Mikel Artetxe, Chungting Zhou, Xi Vic-	(i) any lexicon phrase occurs as a case-	734
687	toria Lin, Ramakanth Pasunuru, Danqi Chen, Luke	insensitive substring, <i>or</i> (ii) any sentence	735
688	Zettlemoyer, and Veselin Stoyanov. 2023. <i>Training</i>	has cosine similarity $\geq \tau$ to a seed exemplar	736
689	<i>trajectories of language models across scales</i> . In	(intfloat/multilingual-e5-small,	737
690	<i>Proceedings of the 61st Annual Meeting of the As-</i>	query:/passage: prefixes), $\tau = 0.82$	738
691	<i>sociation for Computational Linguistics (Volume 1:</i>	primary.	739
692	<i>Long Papers)</i> , pages 13711–13738. Association for		
693	Computational Linguistics.	Lexicon (substring match).	740
694	Mengzhou Xia, Sadhika Malladi, Suchin Gururangan,	• empathy expression (empathy_opener):	741
695	Sanjeev Arora, and Danqi Chen. 2024. <i>LESS: Select-</i>	i'm sorry to hear, i am sorry	742
696	<i>ing influential data for targeted instruction tuning</i> . In	to hear, i'm really sorry, that	743
697	<i>Proceedings of the 41st International Conference on</i>	sounds really, that sounds very,	744
698	<i>Machine Learning</i> .	i can understand, i can imagine,	745
699	Frank Xiao and Santiago Aranguri. 2026. <i>Probe-based</i>	it's understandable, that must	746
700	<i>data attribution: Discovering and mitigating unde-</i>	be, i'm here for you, i hear you,	747
701	<i>sirable behaviors in LLM post-training</i> . Preprint,	it sounds like you	748
702	arXiv:2602.11079. ArXiv preprint, v3.	• validation: your feelings are	749
703	Chungting Zhou, Pengfei Liu, Puxin Xu, Srini Iyer, Jiao	valid, it's okay to feel, it's	750
704	Sun, Yuning Mao, Xuezhe Ma, Avia Efrat, Ping Yu,	normal to feel, it's completely	751
705	Lili Yu, Susan Zhang, Gargi Ghosh, Mike Lewis,	normal, what you're feeling,	752
706	Luke Zettlemoyer, and Omer Levy. 2023. <i>LIMA:</i>	there's nothing wrong with,	753
707	<i>Less is more for alignment</i> . In <i>Advances in Neural</i>	it's natural to	754
708	<i>Information Processing Systems 36 (NeurIPS)</i> .	• professional referral (disclaimer):	755
709	A Prompts and decoding configuration	consult a, talk to a, speak with	756
710	All answers use nucleus sampling (tempera-	a, see a doctor, seek medical,	757
711	ture 0.7, top- p 0.9, max_new_tokens=256,	i am not a doctor, i'm not a	758
712	do_sample=True), the chat template where	doctor, medical professional,	759
713	available (plain Question:/Answer: for base	healthcare professional, mental	760
714	checkpoints), and five seeds per prompt. The bat-	health professional, seek	761
715	ttery: 60 advice, 30 factual, 20 emotional non-	professional, professional help	762
716	advice, 20 neutral advice, 10 domain factual; the		
717	complete lists ship with the artifact.		

- **crisis_referral** (dropped after 0 TP / 11 FP in the 120-sentence audit): crisis, hotline, helpline, 988, call 911, emergency services, suicide
- **structuring language** (structure): here are some, here are a few, consider the following, steps you can take, a few things you can, first,, secondly,

e5 seed exemplars. Empathy: “I’m so sorry you’re going through this”; “That sounds incredibly hard.” Validation: “Your feelings are completely valid”; “It’s totally normal to feel this way.” Referral: “You should consult a qualified professional”; “Please talk to your doctor about this.” Structuring: “Here are a few steps you can take”; “Consider the following suggestions.”

C Threshold sweep

Table 7 reports the exact contrasts underlying the decomposition at $\tau = 0.80/0.82/0.84$. Structuring is stable: advice and neutral-advice remain within .07 in every cell. Validation exceeds its strongest control in five of six cells (the exception is 1B at .80, -.02). Empathy’s emotional-minus-advice contrast reverses at .84, while referral’s advice-minus-domain contrast changes magnitude. The primary $\tau = .82$ was frozen before the 140-prompt battery, but its qualitative decomposition should not be generalized across detector thresholds.

M	Contrast	.80	.82	.84
1B	empathy: emot. – adv.	.013	.203	-.127
7B		.030	.143	-.280
1B	validation: adv. – max ctrl.	-.020	.200	.083
7B		.047	.253	.170
1B	referral: adv. – domain	-.003	.127	.163
7B		.017	.093	.280
1B	structuring: adv. – neutral	.000	.000	.067
7B		.000	-.030	-.010

Table 7: Threshold robustness of the four decomposition contrasts. Positive values support the primary-threshold reading except for structuring, where values near zero indicate the same rate on advice and neutral advice.

D Equivalent-count sensitivity analysis

For readers who want familiar binomial statistics, this table binarizes each prompt’s five-generation mean into an equivalent count ($k = \text{round}(\bar{p} \cdot n)$) and applies Wilson intervals, Fisher exact tests, and

Haldane–Anscombe odds ratios. **Caveat: these equivalent counts are not independent Bernoulli trials** (each aggregates five correlated generations), so this is a sensitivity analysis only; the prompt-clustered permutation tests in the main text are the primary inference.

M	Behavior	Advice	Factual	OR	Fisher p
1B	empathy	35/60	1/30	27.4	1.4×10^{-7}
1B	validation	22/60	0/30	35.6	4.2×10^{-5}
1B	referral	56/60	2/30	143.1	8.9×10^{-17}
1B	structuring	60/60	18/30	81.8	3.2×10^{-7}
7B	empathy	37/60	1/30	31.4	2.2×10^{-8}
7B	validation	23/60	0/30	38.2	1.9×10^{-5}
7B	referral	55/60	2/30	115.0	5.7×10^{-16}
7B	structuring	58/60	16/30	20.6	1.2×10^{-6}

Table 8: Equivalent-count sensitivity analysis (advice vs factual, $\tau = 0.82$). Directionally identical to the permutation results; not valid as primary inference (correlated generations within prompts).

E Per-phrase recoverability and leave-one-out

Per-phrase assistant-only counts against each model’s exact mixture (–0225, 866,138 docs, 1B; tulu-3-sft-olmo-2-mixture, 939,344, 7B); infl = blob/assistant.

Model	Phrase	Asst.	Blob	Infl.
1B	first,	44422	52676	1.19
1B	here are some	11914	13276	1.11
1B	here are a few	5421	6056	1.12
1B	healthcare professional	3075	3216	1.05
1B	mental health professional	2109	2186	1.04
1B	professional help	1765	1825	1.03
1B	seek professional	1353	1377	1.02
1B	it sounds like you	1130	1164	1.03
1B	talk to a	290	410	1.41
1B	it’s natural to	152	157	1.03
1B	it’s completely normal	78	103	1.32
7B	here are some	12049	13432	1.11
7B	here are a few	5511	6156	1.12
7B	healthcare professional	3106	3309	1.07
7B	i’m really sorry	2369	2378	1.00
7B	mental health professional	2120	2214	1.04
7B	professional help	1774	1838	1.04
7B	seek professional	1355	1379	1.02
7B	it sounds like you	1139	1184	1.04
7B	consult a	927	1031	1.11
7B	talk to a	286	408	1.43
7B	it’s completely normal	78	103	1.32

Table 9: Per-phrase SFT-assistant recoverability (representative rows). The 5 (1B) / 6 (7B) distinct 3-word types drive Section 4.3; “here are some” dominates the 3-word mean.

Length-matched buckets (occurrence-weighted). Behavioral vs topical per-million for 1B: 2-word 1839.9 vs 1708.2 ($1.1\times$, $n_{\text{occ}}=11$), 3-word 12333.4 vs 633.8 ($19.5\times$, 46), 4-word 1676.2 vs 11.8 ($141.7\times$, 28). For 7B, 2-word 1784.5 vs 996.8 ($1.8\times$, 20), 3-word 9354.4 vs 1280.4 ($7.3\times$, 55), 4-word 1396.6 vs 96.2 ($14.5\times$, 24). **Distinct-phrase leave-one-out (3-word):** dropping “here are some” moves the mean from 3358.1 to 758.8/M at 1B (5 types) and from 3026.0 to 1065.9/M at 7B (6 types).

F Statistical formulas

Primary inference uses two-sided prompt-clustered permutation tests on prompt-level means (10,000 draws), with Monte Carlo $p = (b + 1)/(B + 1)$; $p < 10^{-4}$ denotes zero raw exceedances. Appendix D reports Wilson 95% intervals, two-sided Fisher tests, and Haldane–Anscombe odds ratios on equivalent counts. The distinct-phrase leave-one-out removes the most frequent 3-word phrase before recomputing the mean behavioral per-million. DPO comparisons use exact McNemar tests on chosen-only versus rejected-only pairs and phrase occurrences per million whitespace-delimited response words. Released scripts derive all statistics from the artifact.